



NETWORK STATISTICS SUCH AS THROUGHPUT, TRANSMISSION SPEED, AVERAGE RTT

A MINI PROJECT

Submitted by

Pragya Chawla – 200905084

Anurag Tiwari - 200905110

In partial fulfilment for the award of the degree of
Bachelor's in technology (B.Tech)

IN

Computer Science & Engineering

BONAFIDE CERTIFICATE

Certified that this project report “**Network Statistics such as Throughput, Transmission Speed, Average RTT.**” is the bonafide work of:

Pragya Chawla (200905084)

&

Anurag Tiwari (200905110)

who carried out the mini project work under my supervision.

Dr. Ashalatha Nayak

Head of Department, CSE

Dr. Roopalakshmi R

Professor, CSE

Submitted to the Viva voce Examination held on

EXAMINER 1

EXAMINER 2

ABSTRACT

For our mini project, we are sniffing data packets passed to us in a .pcap file and calculating different types of statistics regarding each packet. At the end, we analyse the overall statistics of the network such as average speed, average packet size, average packet rate and average RTT. For each packet we try to print the host ip, destination ip, source port address, destination port address, packet capture length, packet total length, sequence number and acknowledgement number. For TCP oriented data packets, we have even printed the payload. The whole code is written in C in which we use the libcap library.

The libpcap library was written as part of a more extensive program called TCPDump. The libpcap library allowed developers to write code to receive link-layer packets (Layer 2 in the OSI model) on different flavours of UNIX operating systems without having to worry about the idiosyncrasy of different operating systems' network cards and drivers. Essentially, the libpcap library grabs packets directly from the network cards, which allowed developers to write programs to decode, display, or log the packets.

In our program, we allow the user to input any .pcap file for which we output its statistics. We also allow the user to put a limit to the number of packets to be sniffed.

ACKNOWLEDGEMENT

We would like to thank our professor, Dr. Roopalakshmi R for providing us with the motivation and the knowledge required to implement this project. He was very helpful and got through to us very clearly regarding all topics on the subject. Once the basics were clear, it was an easy feat to have implemented the project.

We would also like to appreciate the infrastructure provided to us by our educational institute, Manipal Institute of Technology, without which we would have a lot of trouble getting familiar with the required software.

We would also like to acknowledge the help of all our peers and lab staff for helping us understand the problem statement, hence enabling us to be able to write code for the same.

We wish to thank our well-wishers and guardians for their undivided support and interests in our well-being and being a constant source of motivation and inspiration.

TABLE OF CONTENTS

Serial Number	Topic Name	Page Number
1	Cover & Title Page	1
2	Bonafide Certificate	2
3	Abstract	3
4	Acknowledgement	4
5	Table of Contents	5
6	Chapters: a. Network Statistics b. Implementation in Code	6 7
7	Appendices	25
8	References	26

Chapter 1: Network Statistics

1.1 Round Trip Time (RTT)

The round-trip time for a segment is the amount of time between when the segment is sent and when an acknowledgement for the segment is received.

1.2 Throughput

The throughput is the amount of data per second that can be transferred between end systems.

1.3 Processing Delay

The processing delay is the time required to examine the packet's header and determine where to direct the packet.

1.4 Queueing Delay

The queueing delay is the amount of time a packet has to wait before being transmitted onto the link.

1.5 Transmission Delay

The transmission delay is the amount of time required to transmit all of the packet's bits into the link.

1.6 Propagation Delay

The propagation delay is the time required to propagate from the beginning of the link to the end router.

Chapter 2: Implementation in Code

2.1 Introduction

We have developed a command line interface which analyses packets present in a pcap file and present the network statistics such as throughput, average round trip time and transmission speed, for those packets.

This pcap file can downloaded or generated live via any packet sniffing tool such as Wireshark and then be fed as input to our interface.

We have also added a feature giving the user the liberty to decide how many packets the user wants to sniff from the file.

After displaying all packets, its protocols, payload, and respective details, we display the network statistics in the output.

2.2 GitHub Link for the Code

The code has been completed and tested with .pcap and .pcapng files the samples of which along with the source code file has been uploaded in a GitHub repository we have created, the link to which is:

<https://github.com/pragyachawla001/Network-Statistics>

2.3 Walking through the code

2.3.1 Header Files and Libraries used:

The interface was made in the C language. The following header files were used for the implementation:

- `stdio.h`
- `pcap.h`
- `stdlib.h`
- `netinet/in.h`
- `netinet/tcp.h`
- `netinet/udp.h`

- `netinet/ip.h`
- `unistd.h`
- `net/ethernet.h`
- `string.h`

Using libpcap allows us to capture or send packets from a live network device or a file. These code examples will walk us through using libpcap to find network devices, get information about devices, process packets in real time or offline, send packets, and even listen to wireless traffic. This is aimed at Debian based Linux distributions but may also work on Mac OSX. Not intended for Windows, but WinPcap is a port that is available. Compiling a pcap program requires linking with the pcap lib. We can install it in Debian based distributions with:

```
sudo apt-get install libpcap-dev
```

Once the libpcap dependency is installed, we can compile pcap programs with the following command. We will need to run the program as root or with sudo to have permission to access the network card:

```
gcc <filename> -lpcap
```

The code we have written to include the respective header files and libraries used is shown in the screenshot of the snippet of the code below:


```

#include <pcap.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <ctype.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <assert.h>

#define APP_NAME "sniffex"
#define APP_DESC "Sniffer example using libpcap"

/* default snap length (maximum bytes per packet to capture) */
#define SNAP_LEN 1518

/* ethernet headers are always exactly 14 bytes [1] */
#define SIZE_ETHERNET 14

/* Ethernet addresses are 6 bytes */
#define ETHER_ADDR_LEN 6

```

We are including all header files respective to the functions and header structures we have used in our program.

2.3.2 Structures defined for sniffing:

```

/* Ethernet header */
You, 8 hours ago | 1 author (You)
struct sniff_ethernet {
    u_char ether_dhost[ETHER_ADDR_LEN]; /* destination host address */
    u_char ether_shost[ETHER_ADDR_LEN]; /* source host address */
    u_short ether_type; /* IP? ARP? RARP? etc */
};

/* IP header */
You, 8 hours ago | 1 author (You)
struct sniff_ip {
    u_char ip_vhl; /* version << 4 | header length >> 2 */
    u_char ip_tos; /* type of service */
    u_short ip_len; /* total length */
    u_short ip_id; /* identification */
    u_short ip_off; /* fragment offset field */
#define IP_RF 0x8000 /* reserved fragment flag */
#define IP_DF 0x4000 /* don't fragment flag */
#define IP_MF 0x2000 /* more fragments flag */
#define IP_OFFMASK 0x1fff /* mask for fragmenting bits */
    u_char ip_ttl; /* time to live */
    u_char ip_p; /* protocol */
    u_short ip_sum; /* checksum */
    struct in_addr ip_src, ip_dst; /* source and dest address */
};
#define IP_HL(ip) (((ip) -> ip_vhl) & 0x0f)
#define IP_V(ip) (((ip) -> ip_vhl) >> 4)

```

```

/* TCP header */
typedef u_int tcp_seq;

You, 8 hours ago | 1 author (You)
struct sniff_tcp {
    u_short th_sport; /* source port */
    u_short th_dport; /* destination port */
    tcp_seq th_seq; /* sequence number */
    tcp_seq th_ack; /* acknowledgement number */
    u_char th_offx2; /* data offset, rsvd */
#define TH_OFF(th) (((th) -> th_offx2 & 0xf0) >> 4)
    u_char th_flags;
#define TH_FIN 0x01
#define TH_SYN 0x02
#define TH_RST 0x04
#define TH_PUSH 0x08
#define TH_ACK 0x10
#define TH_URG 0x20
#define TH_ECE 0x40
#define TH_CWR 0x80
#define TH_FLAGS (TH_FIN|TH_SYN|TH_RST|TH_ACK|TH_URG|TH_ECE|TH_CWR)
    u_short th_win; /* window */
    u_short th_sum; /* checksum */
    u_short th_urp; /* urgent pointer */
};

You, 8 hours ago | 1 author (You)
struct sniff_udp {
    u_short sport; //source port
    u_short dport; //destination port
    u_short len; //datagram length
    u_short crc; //checksum
};

```

The defined structures and the purpose they serve are mentioned in the comments of the code as shown in the code screenshot above.

The data in structures are relevant to what we read for TCP and UDP headers and packets. Thus, the structures are well defined and explained through the variables used and the comments given in the code.

2.3.3 Prototyping functions required for the code:

```

void got_packet(u_char * args,
    const struct pcap_pkthdr * header,
    const u_char * packet);

void print_payload(const u_char * payload, int len);

void print_hex_ascii_line(const u_char * payload, int len, int offset);

void print_app_usage(void);

```

2.3.4 Functions for printing payload of the packets sniffed:

```
93  /*
94  * print help text
95  */
96  void print_app_usage(void) {
97
98      printf("Usage: %s [interface]\n", APP_NAME);
99      printf("\n");
100     printf("Options:\n");
101     printf("    interface    Listen on <interface> for packets.\n");
102     printf("\n");
103     return;
104 }
105
106 /*
107 * print data in rows of 16 bytes: offset  hex  ascii
108 *
109 * 00000  47 45 54 20 2f 20 48 54  54 50 2f 31 2e 31 0d 0a  GET / HTTP/1.1..
110 */
111 void print_hex_ascii_line(const u_char * payload, int len, int offset) {
112
113     int i;
114     int gap;
115     const u_char * ch;
116
117     /* offset */
118     printf("%05d  ", offset);
119
120     /* hex */
121     ch = payload;
122     for (i = 0; i < len; i++) {
123         printf("%02x ", *ch);
124         ch++;
125         /* print extra space after 8th byte for visual aid */
126         if (i == 7)
127             printf(" ");
128     }
129
130     /* print space to handle line less than 8 bytes */
131     if (len < 8)
132         printf(" ");
133
134     /* fill hex gap with spaces if not full line */
135     if (len < 16) {
136         gap = 16 - len;
137         for (i = 0; i < gap; i++) {
138             printf(" ");
139         }
140     }
141     printf(" ");
142
143     /* ascii (if printable) */
144     ch = payload;
145     for (i = 0; i < len; i++) {
146         if (isprint(*ch))
147             printf("%c", *ch);
148         else
149             printf(".");
150         ch++;
151     }
152
153     printf("\n");
154     return;
155 }
156
157 /*
158 * print packet payload data (avoid printing binary data)
159 */
160 void print_payload(const u_char * payload, int len) {
161
162     int len_rem = len;
163     int line_width = 16; /* number of bytes per line */
164     int line_len;
165     int offset = 0; /* zero-based offset counter */

```

```

166     const u_char * ch = payload;
167
168     if (len <= 0)
169         return;
170
171     /* data fits on one line */
172     if (len <= line_width) {
173         printf(" \t\t");
174         print_hex_ascii_line(ch, len, offset);
175         return;
176     }
177
178     /* data spans multiple lines */
179     while(1) {
180         /* compute current line length */
181         line_len = line_width % len_rem;
182         /* print line */
183         printf(" \t\t");
184         print_hex_ascii_line(ch, line_len, offset);
185         /* compute total remaining */
186         len_rem = len_rem - line_len;
187         /* shift pointer to remaining bytes to print */
188         ch = ch + line_len;
189         /* add offset */
190         offset = offset + line_width;
191         /* check if we have line width chars or less */
192         if (len_rem <= line_width) {
193             /* print last line and get out */
194             printf(" \t\t");
195             print_hex_ascii_line(ch, len_rem, offset);
196             break;
197         }
198     }
199
200     return;
201 }

```

The functions “print_payload()” and “print_hex_ascii_line()” are both functions which are used to help print the payload of the packets which we are sniffing from the pcap file fed as input to the program. These files use the header information and other formatting details to produce the payload which we then print individually for each packet sniffed.

2.3.5 Function for sniffing and displaying each packet individually:

got_packet is a function used as the call-back for our pcap_loop function. It takes in three arguments, a u_char pointer which is passed in the user argument to pcap_loop(), a const struct pcap_pkthdr pointer pointing to the packet timestamp and lengths, and a const u_char pointer to the first caplen (as given in the struct pcap_pkthdr a pointer to which is passed to the callback routine) bytes of data from the packet. The struct pcap_pkthdr and the packet data are not to be freed by the callback routine, and are not guaranteed to be valid after the callback routine returns; if the code needs them to be valid after the callback, it must make a copy of them. Count is used for keeping a count of packets. ethernet is a pointer to sniff_ethernet struct, which points to the

starting of the packet. ip is a pointer to sniff_ip struct, which points to the start of the packet with an offset equal to the size of the ethernet header. tcp is a pointer to sniff_tcp which points to the start of the packet with an offset of ethernet size, ip size combined. We then get to know which type of protocol is being used with the help of the switch case. Our program will continue only if the packet follows TCP or UDP protocol. For UDP, we print the source port number, destination port number and datagram length. For TCP, we print the source port number, destination port number, packet capture length, packet total length, sequence number, acknowledgement number and payload, if any.

```

203  /*
204  * dissect/print packet
205  */
206  void got_packet(u_char * args,
207      const struct pcap_pkthdr * header,
208      const u_char * packet) {
209
210      static int count = 1; /* packet counter */
211
212      /* declare pointers to packet headers */
213      const struct sniff_ethernet * ethernet; /* The ethernet header [1] */
214      const struct sniff_ip * ip; /* The IP header */
215      const struct sniff_tcp * tcp; /* The TCP header */
216      const char * payload; /* Packet payload */
217
218      int size_ip;
219      int size_tcp;
220      int size_payload;
221
222      printf("\t\t----- PACKET [Number : %d] ----- \n\n", count++);
223
224      /* define ethernet header */
225      ethernet = (struct sniff_ethernet *) (packet);
226
227      /* define/compute ip header offset */
228      ip = (struct sniff_ip *) (packet + SIZE_ETHERNET);
229      size_ip = IP_HL(ip) * 4;
230      if (size_ip < 20) {
231          printf("***\t\tInvalid IP header length: %u bytes\n", size_ip);
232          return;
233      }
234
235      /* print source and destination IP addresses */
236      printf("***\t\tFrom IP : %s\n", inet_ntoa(ip -> ip_src));
237      printf("***\t\tTo IP : %s\n", inet_ntoa(ip -> ip_dst));

```

```

238
239 /* determine protocol */
240 switch (ip -> ip_p) {
241 case IPPROTO_TCP:
242     printf("***\t\tProtocol : TCP\n");
243     break;
244 case IPPROTO_UDP:
245     printf("***\t\tProtocol : UDP\n");
246     break;
247 case IPPROTO_ICMP:
248     printf("***\t\tProtocol : ICMP\n");
249     return;
250 case IPPROTO_IP:
251     printf("***\t\tProtocol : IP\n");
252     return;
253 default:
254     printf("***\t\tProtocol : Unknown\n");
255     return;
256 }
257
258 /* define/compute tcp header offset */
259 if (ip -> ip_p == IPPROTO_UDP) {
260     struct sniff_udp * udp;
261     udp = (struct sniff_udp *) (packet + SIZE_ETHERNET + size_ip);
262     printf("***\t\tSource Port : %u\n**\t\tDestination Port : %u\n", udp -> sport, udp -> dport);
263     printf("***\t\tUDP Datagram Length : %u\n", udp -> len / 256);
264 } else {
265
266     tcp = (struct sniff_tcp *) (packet + SIZE_ETHERNET + size_ip);
267     size_tcp = TH_OFF(tcp) * 4;
268     if (size_tcp < 20) {
269         printf("***\t\tInvalid TCP header length: %u bytes\n", size_tcp);
270         return;
271     }

```

```

272
273     printf("***\t\tSource port : %d\n", ntohs(tcp -> th_sport));
274     printf("***\t\tDestination port : %d\n", ntohs(tcp -> th_dport));
275
276     printf("***\t\tPacket capture length : %d\n", header -> caplen);
277     printf("***\t\tPacket total length : %d\n", header -> len);
278
279     printf("***\t\tSequence number : %u\n", tcp -> th_seq);
280     printf("***\t\tAcknowledgement number(Ack) : %u\n", tcp -> th_ack);
281
282     /* define/compute tcp payload (segment) offset */
283     payload = (u_char *) (packet + SIZE_ETHERNET + size_ip + size_tcp);
284
285     /* compute tcp payload (segment) size */
286     size_payload = ntohs(ip -> ip_len) - (size_ip + size_tcp);
287
288     /*
289     * Print payload data; it might be binary, so don't just
290     * treat it as a string.
291     */
292     if (size_payload > 0) {
293         printf("***\t\tPayload (%d bytes) : \n\n", size_payload);
294         print_payload(payload, size_payload);
295     }
296 }
297 printf("\n\t\t-----\n\n");
298 return;
299 }
300

```

2.3.6 Function for parsing the file to give final network statistics output:

We use this function to Parse our data.txt file and get all the valuable information from it. FILE* f points to the file which we want to open (data.txt). buf array is used to store the whole line of our input, whereas temp array is used to store the value of our data, such as average speed, average packet size, average packet rate, etc. If we open the data.txt file, we will soon come to know that the eleventh and twelfth lines of the file hold the average speed of our network, the thirteenth line holds the average packet size, and the fourteenth line holds the average packet rate. The count variable is used to calculate the line number we have parsed so far. Unit conversions are done accordingly to have a standardized way of presenting our data. In the end, we calculate the average RTT of our network and output it accordingly.

Function: parse()

```
300
301 void parse() {
302     FILE * f;
303     f = fopen("data.txt", "r");
304     char buf[256];
305     char temp[256];
306     int count = 0;
307     int i = 0, j = 0;
308     float size, speed, Mbps, prate;
309     for (int count = 0; fgets(buf, sizeof(buf), f) != NULL && count < 15; count++) {
310         if (count < 11) {
311             continue;
312         }
313         i = 0, j = 0;
314         while (!isdigit(buf[i])) i++;
315         while (buf[i] != ' ') {
316             if (buf[i] != ',') {
317                 temp[j++] = buf[i];
318             }
319             i++;
320         }
321         temp[j++] = 0;
322
323         if (count == 11) {
324             speed = atof(temp);
325
326             while(buf[i++] != ' ');
327             char type[10];
328             int k = 0;
329             while(isalpha(buf[i]) || buf[i] == '/') {
330                 type[k++] = buf[i++];
331             }
332             type[k++] = 0;
333             if(strcmp(type, "kBps") == 0) speed /= 1024.0f;
```

```

334     else if(strcmp(type, "bytes/s") == 0) speed /= 1024.0f * 1024.0f;
335     else if(strcmp(type, "Mbps") != 0) puts(type), assert(0);
336
337     printf("**\t\tAVERAGE SPEED(Mbps)    : %.2f Mbps\n", speed);
338 } else if (count == 12) {
339     Mbps = atof(temp);
340
341     while(buf[i++] != ' ');
342     char type[10];
343     int k = 0;
344     while(isalpha(buf[i])){
345         type[k++] = buf[i++];
346     }
347     type[k++] = 0;
348     if(strcmp(type, "kpbs") == 0) Mbps /= 1024.0f;
349     else if(strcmp(type, "Mbps") != 0) assert(0);
350
351     printf("**\t\tAVERAGE SPEED(Mbps)    : %.2f Mbps\n", Mbps);
352 } else if (count == 13) {
353     size = atof(temp);
354     printf("**\t\tAVERAGE PACKET SIZE    : %.2f bytes\n", size);
355 } else if (count == 14) {
356     prate = atof(temp);
357     printf("**\t\tAVERAGE PACKET RATE/s : %.2f kpackets/s\n", prate);
358 }
359 }
360 printf("**\t\tAVERAGE RTT                : %f seconds\n", size * 2 / (speed * (1 << 20)));
361 }
362

```

2.3.7 Function for displaying simple terminal-based GUI:

```

362
363 void menu() {
364     system("clear");
365     printf("\n***** MENU *****");
366     printf("\n\n**      Enter the file name with the .pcap extension for analysis      **");
367     printf("\n\n**      ==>      ");
368 }
369

```

We have made a small terminal-based interface and to beautify, we have used this function and given proper spacing.

2.3.8 The main function that ties all the program together:

To obtain the bare-bones statistics of the network, we used the capinfos software, and stored the fetched data into a text file called “data.txt”. We have then sent this file to be parsed and statistics to be retrieved through the parse() function call mentioned in the code.

We also obtained the IP header information and the TCP/UDP header information.


```

370 int main(int argc, char ** argv) {
371     char errbuf[PCAP_ERRBUF_SIZE]; /* error Temporary */
372     pcap_t * handle; /* packet capture handle */
373
374     char filter_exp[] = "ip"; /* filter expression [3] */
375     struct bpf_program fp; /* compiled filter program (expression) */
376     bpf_u_int32 mask = 0; /* subnet mask */
377     bpf_u_int32 net = 0; /* ip */
378     int num_packets = 0; /* number of packets to capture */
379
380     char fileName[100];
381     menu();
382     scanf("%s", fileName);
383     handle = pcap_open_offline(fileName, errbuf);
384     if (handle == NULL) {
385         fprintf(stderr, "Couldn't open device\n");
386         exit(EXIT_FAILURE);
387     }
388     printf("\n*****");
389     printf("\n\n**      Enter number of packets to be sniffed (Enter 0 for all):      **");
390     printf("\n\n**      ==>      ");
391     scanf("%d", & num_packets);
392     printf("\n*****");
393     if (num_packets == 0)
394         printf("\n\n**      Number of packets : All\n");
395     else
396         printf("\n\n**      Number of packets : %d\n", num_packets);
397
398     /* make sure we're capturing on an Ethernet device [2] */
399     if (pcap_datalink(handle) != DLT_EN10MB) {
400         fprintf(stderr, "Not an Ethernet\n");
401         exit(EXIT_FAILURE);
402     }

```

```

403     printf("\n*****\n\n");
404
405     /* compile the filter expression */
406     if (pcap_compile(handle, & fp, filter_exp, 0, net) == -1) {
407         fprintf(stderr, "Couldn't parse filter %s: %s\n",
408             filter_exp, pcap_geterr(handle));
409         exit(EXIT_FAILURE);
410     }
411
412     /* apply the compiled filter */
413     if (pcap_setfilter(handle, & fp) == -1) {
414         fprintf(stderr, "Couldn't install filter %s: %s\n",
415             filter_exp, pcap_geterr(handle));
416         exit(EXIT_FAILURE);
417     }
418
419     /* now we can set our callback function */
420     pcap_loop(handle, num_packets, got_packet, NULL);
421
422     /* cleanup */
423     pcap_freecode(& fp);
424     pcap_close(handle);
425
426     char command[200];
427     sprintf(command, "capinfos %s > ./data.txt", fileName);
428     system(command);
429     printf("\n\n***** FINAL NETWORK STATISTICS *****\n\n");
430     parse();
431     printf("\n\n*****\n\n");
432
433     printf("\nCapture complete.\n\n");
434
435     return 0;
436 }
437

```

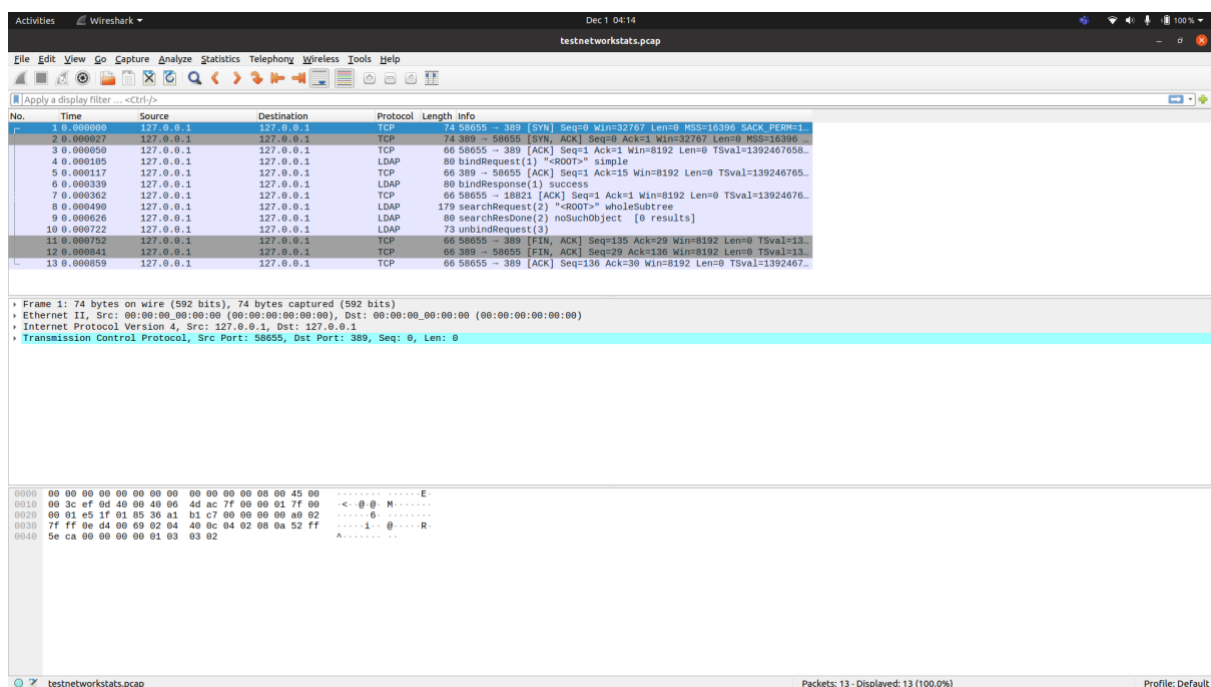
We have used some print statements to beautify the code in the end to compensate for a functioning frontend.

The overall code works and has been tested with multiple pcap and pcapng files and the order of working is that initially on running the code, we are asked to input the name of the pcap file to be fed as input. On inputting a valid file name, we can enter how many packets we want the program to sniff, and we get the appropriate display as chosen. In the end, after all the individual packets have been displayed, we can see that the network statistics are displayed.

We have verified the correctness of the statistics by checking network statistics in the Wireshark tool as well.

2.3.9 Outputs of the Program:

The first sample .pcap file is called “testnetworkstats.pcap” and the Wireshark screenshot is attached below:



We can see the packets available to be sniffed in the screenshot above and now we will run the code in the gcc compiler.

The menu driven GUI looks as follows:

We are required to input the Filename and the number of packets we want to be sniffed.

```
***** MENU *****

**      Enter the file name with the .pcap extension for analysis      **
**      ==>  testnetworkstats.pcap                                     **

*****

**      Enter number of packets to be sniffed (Enter 0 for all):      **
**      ==>  0                                                         **

*****

**      Number of packets : All                                       **

*****

          ----- PACKET [Number : 1] -----

**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 389
**      Packet capture length : 74
**      Packet total length : 74
**      Sequence number : 3350307126
**      Acknowledgement number(Ack) : 0

          -----

          ----- PACKET [Number : 2] -----

**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 389
**      Destination port : 58655
**      Packet capture length : 74
**      Packet total length : 74
**      Sequence number : 750250550
**      Acknowledgement number(Ack) : 3367084342
```

```

-----
----- PACKET [Number : 3] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 389
**      Packet capture length : 66
**      Packet total length : 66
**      Sequence number : 3367084342
**      Acknowledgement number(Ack) : 767027766
-----

----- PACKET [Number : 4] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 389
**      Packet capture length : 80
**      Packet total length : 80
**      Sequence number : 3367084342
**      Acknowledgement number(Ack) : 767027766
**      Payload (14 bytes) :

00000  30 0c 02 01 01 60 07 02  01 03 04 00 80 00      0....`.....
-----

----- PACKET [Number : 5] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 389
**      Destination port : 58655
**      Packet capture length : 66
**      Packet total length : 66
**      Sequence number : 767027766
**      Acknowledgement number(Ack) : 3601965366

```

```

-----
----- PACKET [Number : 6] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 389
**      Destination port : 58655
**      Packet capture length : 80
**      Packet total length : 80
**      Sequence number : 767027766
**      Acknowledgement number(Ack) : 3601965366
**      Payload (14 bytes) :

00000  30 0c 02 01 01 61 07 0a  01 00 04 00 04 00      0....a.....
-----

----- PACKET [Number : 7] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 18821
**      Packet capture length : 66
**      Packet total length : 66
**      Sequence number : 3601965366
**      Acknowledgement number(Ack) : 1001908790
-----

----- PACKET [Number : 8] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 389
**      Packet capture length : 179
**      Packet total length : 179
**      Sequence number : 3601965366
**      Acknowledgement number(Ack) : 1001908790
**      Payload (113 bytes) :

```

```

-----
----- PACKET [Number : 8] -----
**
** From IP : 127.0.0.1
** To IP : 127.0.0.1
** Protocol : TCP
** Source port : 58655
** Destination port : 389
** Packet capture length : 179
** Packet total length : 179
** Sequence number : 3601965366
** Acknowledgement number(Ack) : 1001908790
** Payload (113 bytes) :
00000 30 6f 02 01 02 63 6a 04 00 0a 01 02 0a 01 00 02 00...cj.....
00016 01 00 02 01 00 01 01 00 a9 3c 81 1c 32 2e 31 36 .....<..2.16
00032 2e 38 34 30 2e 07 2e 31 31 33 37 33 30 2e 33 2e .840...113730.3.
00048 33 2e 32 2e 34 36 2e 31 82 10 64 65 70 61 72 74 3.2.46.1..depart
00064 6d 65 6e 74 4e 75 6d 62 65 72 83 07 3e 3d 4e 34 mentNumber...>=N4
00080 37 30 39 84 01 ff 30 19 04 02 63 6e 04 02 73 6e 709...0...cn..sn
00096 04 0f 74 65 6c 65 70 68 6f 6e 65 4e 75 6d 62 65 ..telephoneNumbe
00112 72 r

-----
----- PACKET [Number : 9] -----
**
** From IP : 127.0.0.1
** To IP : 127.0.0.1
** Protocol : TCP
** Source port : 389
** Destination port : 58655
** Packet capture length : 80
** Packet total length : 80
** Sequence number : 1001908790
** Acknowledgement number(Ack) : 1202889014
** Payload (14 bytes) :
00000 30 0c 02 01 02 65 07 0a 01 20 04 00 04 00 0....e... ..
-----

```

```

-----
----- PACKET [Number : 10] -----
**
** From IP : 127.0.0.1
** To IP : 127.0.0.1
** Protocol : TCP
** Source port : 58655
** Destination port : 389
** Packet capture length : 73
** Packet total length : 73
** Sequence number : 1202889014
** Acknowledgement number(Ack) : 1236789814
** Payload (7 bytes) :
00000 30 05 02 01 03 42 00 0....B.

-----
----- PACKET [Number : 11] -----
**
** From IP : 127.0.0.1
** To IP : 127.0.0.1
** Protocol : TCP
** Source port : 58655
** Destination port : 389
** Packet capture length : 66
** Packet total length : 66
** Sequence number : 1320329526
** Acknowledgement number(Ack) : 1236789814

-----
----- PACKET [Number : 12] -----
**
** From IP : 127.0.0.1
** To IP : 127.0.0.1
** Protocol : TCP
** Source port : 389
** Destination port : 58655
** Packet capture length : 66
** Packet total length : 66
** Sequence number : 1236789814
** Acknowledgement number(Ack) : 1337106742

```

```

-----
----- PACKET [Number : 13] -----
**      From IP : 127.0.0.1
**      To IP : 127.0.0.1
**      Protocol : TCP
**      Source port : 58655
**      Destination port : 389
**      Packet capture length : 66
**      Packet total length : 66
**      Sequence number : 1337106742
**      Acknowledgement number(Ack) : 1253567030
-----

***** FINAL NETWORK STATISTICS *****

**      AVERAGE SPEED(MBps)      : 1.18 MBps
**      AVERAGE SPEED(Mbps)      : 9.42 Mbps
**      AVERAGE PACKET SIZE      : 79.69 bytes
**      AVERAGE PACKET RATE/s    : 15.00 kpackets/s
**      AVERAGE RTT              : 0.000129 seconds

*****

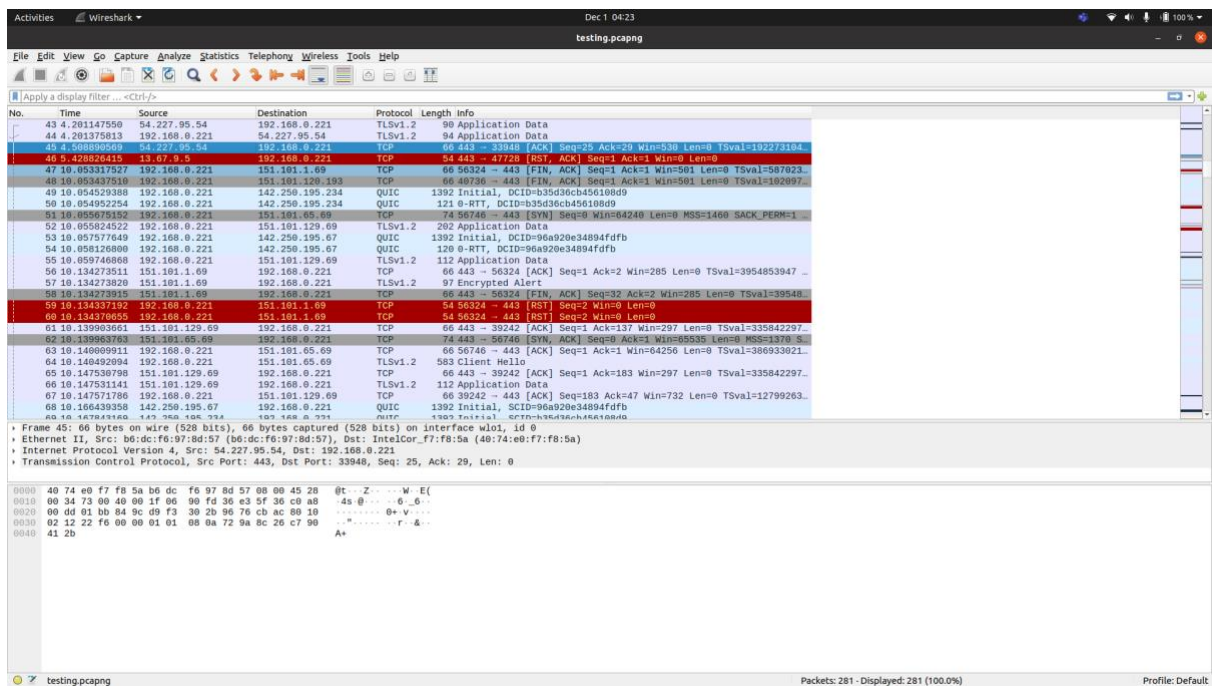
Capture complete.

```

In this final screenshot, we can notice that after all thirteen packets are displayed, we are able to see the network statistics displayed which is the outcome required for our project.

To test our final project, we have used another sample input file called “testing.pcapng” which is very huge in comparison, and we can see the packets and network statistics for that too as shown below:

Initially, we show the Wireshark screenshot for the given pcap file:



The menu driven inputs:

```
***** MENU *****

**      Enter the file name with the .pcap extension for analysis      **

**      ==>   testing.pcapng                                           **

*****

**      Enter number of packets to be sniffed (Enter 0 for all):      **

**      ==>   0                                                         **

*****

**      Number of packets : All                                       **

*****

----- PACKET [Number : 1] -----

**      From IP : 157.240.192.52
**      To IP : 192.168.0.221
**      Protocol : TCP
**      Source port : 443
**      Destination port : 36576
**      Packet capture length : 99
**      Packet total length : 99
**      Sequence number : 1953468210
**      Acknowledgement number(Ack) : 2926838274
**      Payload (33 bytes) :

0000    17 03 03 00 1c 4c ae 01 dd 40 98 bd dc 32 83 d4    ....L...@...2..
0001    42 f6 ba 43 e6 a1 16 c3 c2 08 9e 9d 01 ea 9c fa    B..C.....
0002    8b                                              .

-----
```

The last packet along with the network statistics as shown below:

```
----- PACKET [Number : 279] -----
**      From IP : 54.227.95.54
**      To IP : 192.168.0.221
**      Protocol : TCP
**      Source port : 443
**      Destination port : 33948
**      Packet capture length : 66
**      Packet total length : 66
**      Sequence number : 1127281625
**      Acknowledgement number(Ack) : 3368777366
-----

***** FINAL NETWORK STATISTICS *****
**      AVERAGE SPEED(MBps)      : 0.01 MBps
**      AVERAGE SPEED(Mbps)      : 0.06 Mbps
**      AVERAGE PACKET SIZE      : 465.27 bytes
**      AVERAGE PACKET RATE/s    : 17.00 kpackets/s
**      AVERAGE RTT              : 0.112262 seconds

*****

Capture complete.
```


APPENDICES

GitHub link to our code and input test files:

<https://github.com/pragyachawla001/Network-Statistics>

Link for downloading sample pcap and pcapng files of all sizes and formats along with protocol filters required:

<https://www.wireshark.org/download/automated/captures/>

REFERENCES

1. Computer Networking: A Top-Down Approach, 6th Edition
2. <https://www.opensourceforu.com/2011/02/capturing-packets-c-program-libpcap/>
3. <https://www.devdungeon.com/content/using-libpcap-c>
4. <http://yuba.stanford.edu/~casado/pcap/section4.html>
5. <https://www.wireshark.org/docs/man-pages/capinfos.html>
6. <https://danielmiessler.com/study/tcpdump/>